



MongoDB

humongous database



En bref

- Developed par : MongoDB Inc., California
- Under licence : AGPL
- Current version : 2.6
- Written in C++



SQL VS noSQL

SQL	noSQL
Relational BDD	No FK, no joins
Based on tables	Based on collections, key-value pairs, graphs
Based on SQL language	Based on Unstructured Query Language, defined by each systems
Good for complex requests	Bad for complex requests
Vertically scalable	Horizontally scalable
Good for transactional purposes	Does not support transactions



Quick analogy with RDBMS

table ↔ collection

row ↔ BSON document

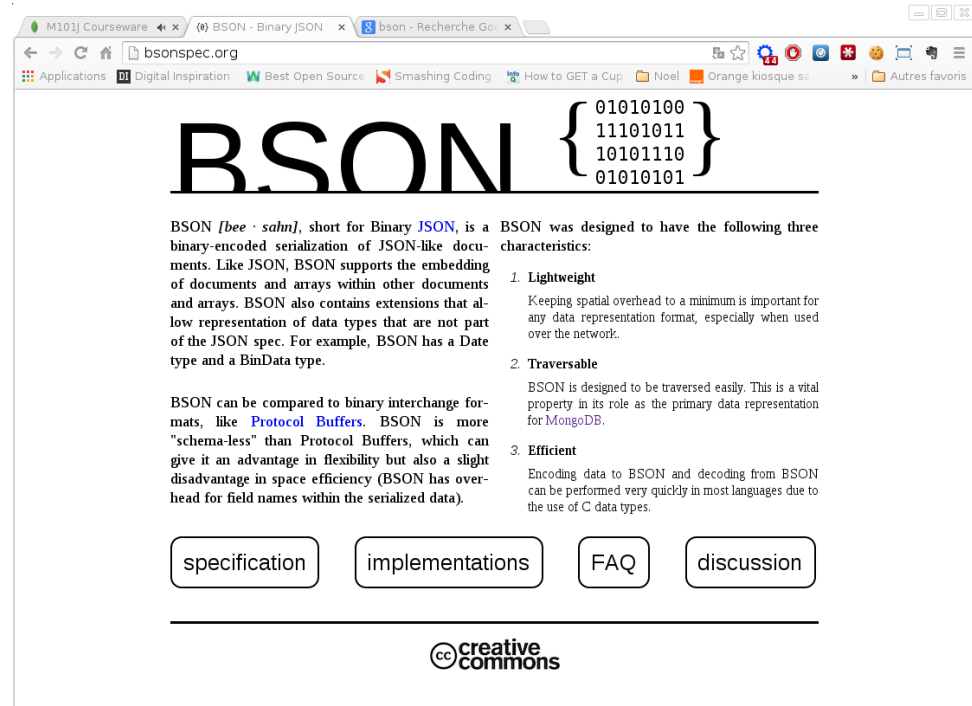

shard ↔ instance

BSON ?

BSON

=

superset(JSON)



Extended types : utf8 strings, binary type, datetime, regular expressions, js code, double precision float, etc



CRUD

MongoDB provides a javascript shell to perform CRUD operations and to configure database

insert : db.collection.**insert**(doc)

select : db.collection.**find**()

where : db.collection.**find**(*conditions*)

update : db.collection.**update**(conditions, modif)

Operators : \$gt, \$or, \$and, \$in, \$all, \$set, \$push, \$pull

```
db.users.find( { friends : { $all : [ "Joe" , "Bob" ] },  
favorites : { $in : [ "running" , "pickles" ] } } )
```



MongoDB cursor

In mongoDB, the result of a request is a cursor

```
cur = db.users.find();  
cur.hasNext();  
cur.next();  
cur.sort({name : -1});  
cur.sort({name : -1}).skip(30).limit(3);  
cur.count();
```



Schema design

- Schemaless but mostly you use a schema
- Application driven schema
- Pre join/Embed data
- No constraints

NoSQL rule :

“Match the data access patterns of your application”




Schema design

```
{ "_id" : "3266F86FD7FD",  
  "postContent" : "J'aime les bananes",  
  "comments" : [  
    { author : "Jacky", "email": "j.terin@gmail.com" },  
    { author : "Mich", "email": "m.terin@gmail.com" },  
  ],  
  tags : [ "eating", "fruits", "yellow" ]  
}
```

One to many relationship

people

```
{ "_id" : "3266F86FD7FD",  
  "name" : "Patrick",  
  "city" : "NYC" }
```



cities

```
{ "_id" : "NYC",  
  ... }
```

But embedded
arrays should
be preferred
when the
many is not
large



Many to many relationship

people

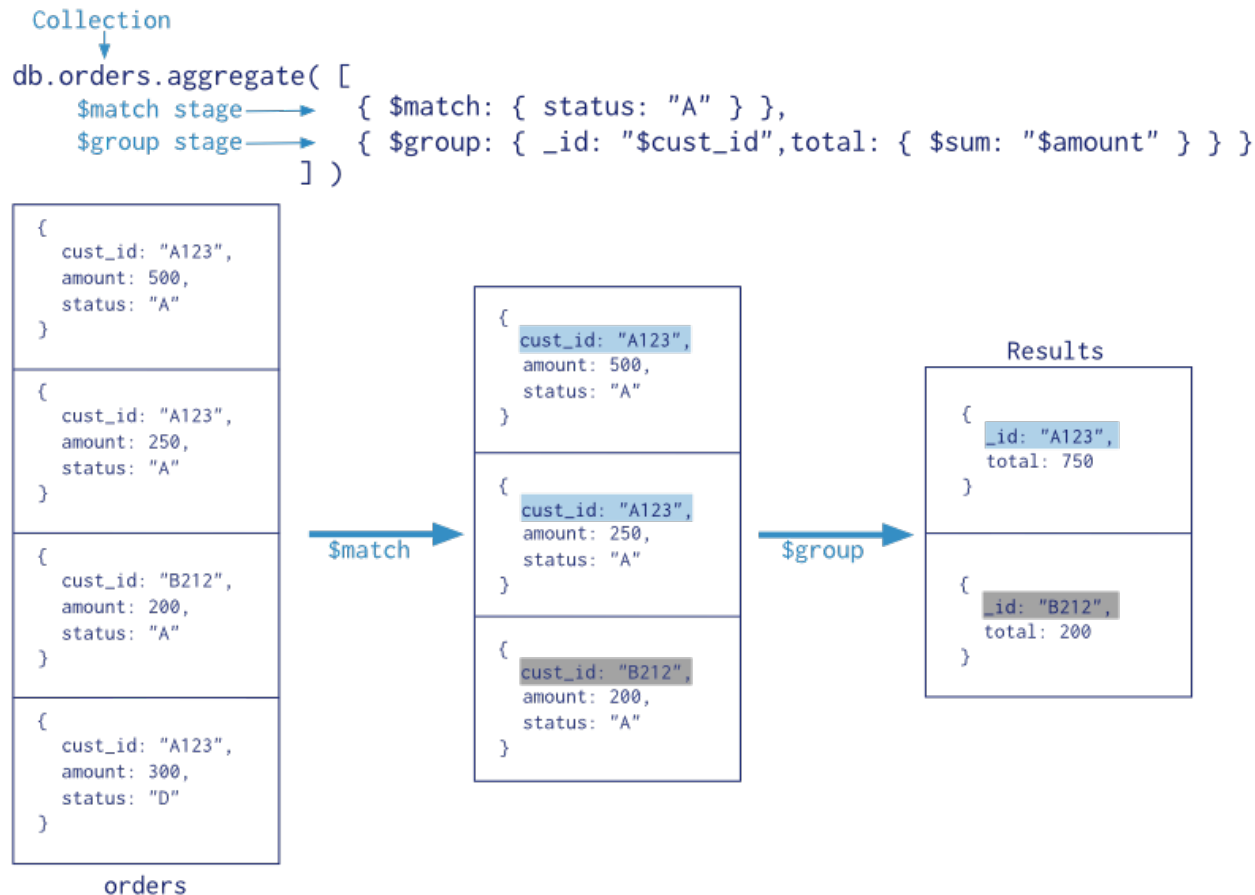
```
{ "_id" : "3266F86FD7FD",  
  "name" : "Patrick",  
  "lived_in" : [12, 5, 71] }
```

cities

```
{ "_id" : 12,  
  "name" : "New York City"  
  ... }
```


Aggregation in a few words

Aggregations process data records and return computed results



Performance - Indexes

Indexes goal : Reduce database scanning



```
db.collection.ensureIndex({user_name:1})  
db.collection.getIndexes()
```



Performance - Explain function

`db.collection.find(conditions).explain()`

```
{  
  "cursor" : "<Cursor Type and Index>",  
  "isMultiKey" : <boolean>,  
  "n" : <num>,  
  "nscannedObjects" : <num>,  
  "nscanned" : <num>,  
  "nscannedObjectsAllPlans" : <num>,  
  "nscannedAllPlans" : <num>,  
  "scanAndOrder" : <boolean>,  
  "indexOnly" : <boolean>,  
  "nYields" : <num>,  
  "nChunkSkips" : <num>,  
  "millis" : <num>,  
  "indexBounds" : { <index bounds> },  
}
```




Performance - MongoDB Profiler

- Start your mongod process with profile command
- Allow you to save all the queries, or queries that takes longer than X ms, where X is a parameter
- Insert queries information in a collection called system.profile

`db.system.profile.find()`

Application Engineering - Replication

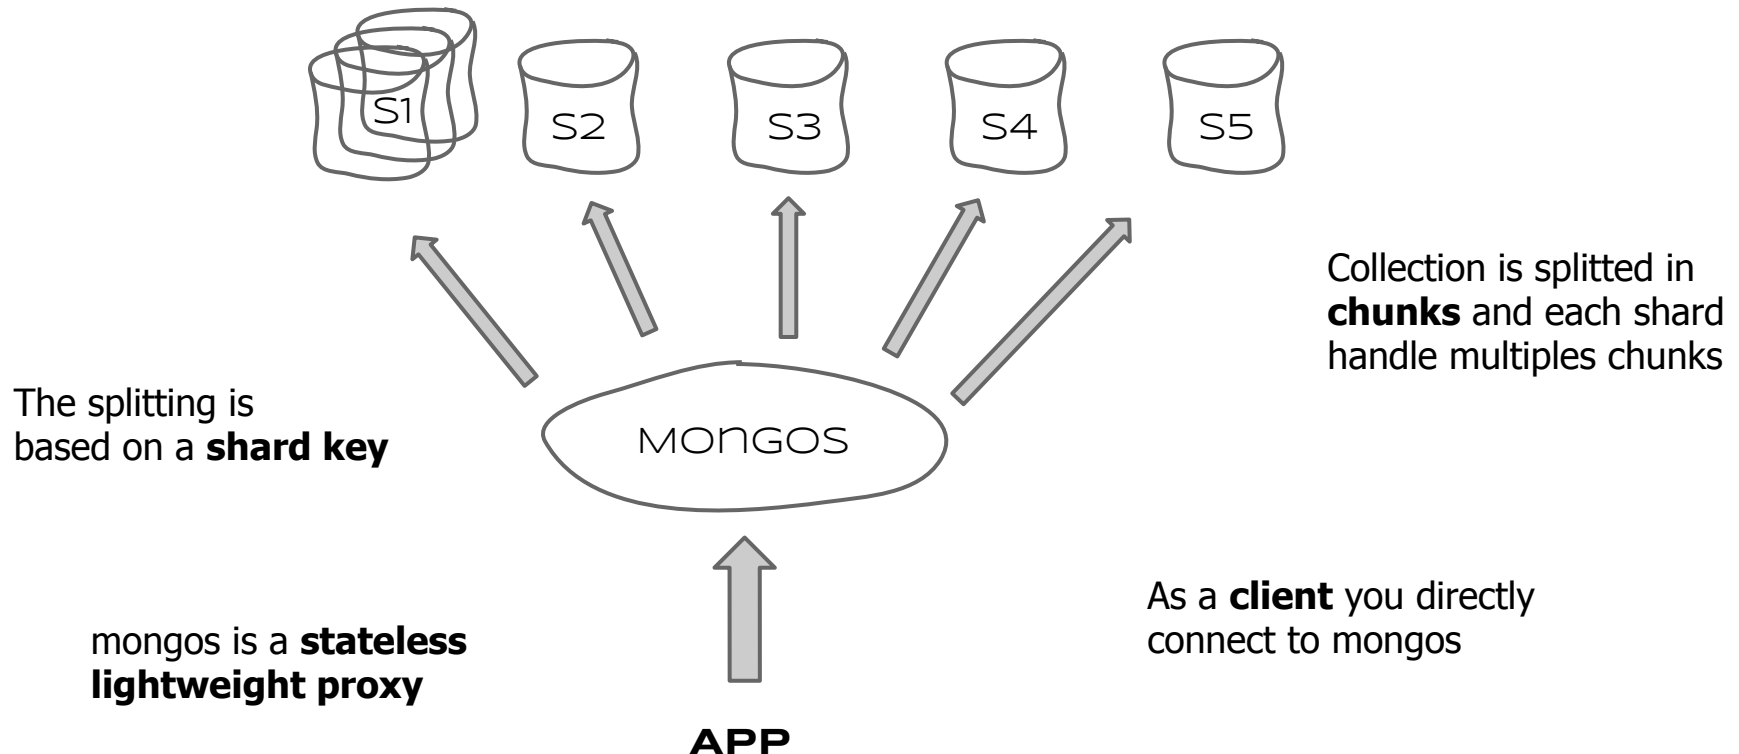
Replication needed for Availability and Fault tolerance



Three nodes are needed to assure the election of a new Primary

Application Engineering - Sharding

- Sharding is the way you can scale out
- Data is distributed on multiple servers





Application Engineering - Sharding

- How to choose a shard key ?
 - Sufficient cardinality for a good distribution
 - Avoid hotspotting (ie, a shardkey that always increase)
- Who set the chunk ranges ?
 - By default, mongos create chunks based on shardkey type (ie, integer from 1 to 2^{32})
 - Chunks can be defined manually
- And if the mongos is down ?
 - Actually the mongos has to be replicated too

NoSQL comparison

Features



elasticsearch.

mongoDB

SQLite

APACHE
HBASE

redis

cassandra

Performances



Online test

goo.gl/Xs2zMR

Start working!